

End-to-End ML Pipeline

Imagine you're baking a cake. You gather ingredients (data ingestion), mix and prepare them (data cleaning, feature engineering), bake the cake (train your model), taste-test it (evaluate), then serve it to others (deploy and monitor). That's an ML pipeline — taking data from raw to useful predictions.

Supervised, Unsupervised & Reinforcement Learning

Supervised: You give the computer the question and the answer, and it learns to match them (like teaching with flashcards). Unsupervised: You just give the computer the data, and it finds patterns on its own (like sorting LEGO pieces by shape without labels). Reinforcement: The computer learns by trying, getting rewards for good actions (like a dog learning tricks with treats).

Time Series Forecasting & Monte Carlo Simulation

Time series is like trying to guess tomorrow's weather based on past patterns. Monte Carlo simulation is like rolling dice many times to guess the chances of something happening — helpful when you're unsure and want to understand all possible outcomes.

NLP & LLMs

NLP is how computers read, understand, and write human language. LLMs (like ChatGPT) are trained on lots of text and can answer questions, write summaries, or even draft emails — like a very smart intern who reads a million books a day.

MLOps + Cloud Deployment

MLOps is like DevOps for ML — it helps you build, ship, and maintain ML models smoothly. Cloud tools like AWS let you run projects on someone else's supercomputer, so you can scale models reliably, like turning a food truck into a restaurant chain.

Data Visualization & Storytelling

Data visualization turns numbers into pictures so people can see patterns — think graphs, charts, dashboards. Storytelling is explaining those patterns so someone can make a smart decision — like turning a traffic map into 'Take Route B — it's faster!'

Real-World Data (RWD)

RWD is messy, like kids' drawings on paper — lots of creativity, but not always clean. It comes from real people: hospital records, prescriptions, sales logs. You have to clean it, match things up, and make sure it's trustworthy before using it.

Version Control & Collaboration (Git, Agile)

Git is like Google Docs for code — it tracks changes, lets you go back, and lets your team work together. Agile means breaking big jobs into small pieces, sharing updates often, and adjusting as you go — like building LEGO castles one section at a time.